

## CIND 820- Big Data Analytics Project

### Using Machine Learning for Prediction of Early Readmission of Diabetic Patients

Supervised by: Ceni Babaoglu

Presented by: Nehal Gamal Mohamed (501278190)

```
In [ ]: !pip install pandas
import sys
!pip install matplotlib
!pip install graphviz
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.0.3)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.1)
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/dist-packages (from pandas) (1.25.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.2.1)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.53.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.25.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (2.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-packages (0.20.3)
```

```
In [ ]: #importing necessary libraries
import csv
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
In [ ]: #Uploading diabetic_data csv file
from google.colab import files
data = files.upload()
```

```
#Instintiating DataReader class
class DataReader:
    def read_csv(self, filename):
        df = pd.read_csv(filename)
        return df

filename = next(iter(data))
data_reader = DataReader()
df = data_reader.read_csv(filename)
#Displaying the first 10 records of the dataset
print(df.head(10))
```

Choose Files

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving diabetic\_data (version 1).csv to diabetic\_data (version 1) (1).csv

	encounter_id	patient_nbr	race	gender	age	weight	\
0	2278392	8222157	Caucasian	Female	[0-10)	?	
1	149190	55629189	Caucasian	Female	[10-20)	?	
2	64410	86047875	AfricanAmerican	Female	[20-30)	?	
3	500364	82442376	Caucasian	Male	[30-40)	?	
4	16680	42519267	Caucasian	Male	[40-50)	?	
5	35754	82637451	Caucasian	Male	[50-60)	?	
6	55842	84259809	Caucasian	Male	[60-70)	?	
7	63768	114882984	Caucasian	Male	[70-80)	?	
8	12522	48330783	Caucasian	Female	[80-90)	?	
9	15738	63555939	Caucasian	Female	[90-100)	?	

	admission_type_id	discharge_disposition_id	admission_source_id	\
0	6	25	1	
1	1	1	7	
2	1	1	7	
3	1	1	7	
4	1	1	7	
5	2	1	2	
6	3	1	2	
7	1	1	7	
8	2	1	4	
9	3	3	4	

	time_in_hospital	...	citoglipton	insulin	glyburide-metformin	\
0	1	...	No	No	No	
1	3	...	No	Up	No	
2	2	...	No	No	No	
3	2	...	No	Up	No	
4	1	...	No	Steady	No	
5	3	...	No	Steady	No	
6	4	...	No	Steady	No	
7	5	...	No	No	No	
8	13	...	No	Steady	No	
9	12	...	No	Steady	No	

	glipizide-metformin	glimepiride-pioglitazone	metformin-rosiglitazone	\
0	No		No	No
1	No		No	No
2	No		No	No
3	No		No	No
4	No		No	No
5	No		No	No
6	No		No	No
7	No		No	No
8	No		No	No
9	No		No	No

	metformin-pioglitazone	change	diabetesMed	readmitted
0	No	No	No	NO
1	No	Ch	Yes	>30
2	No	No	Yes	NO
3	No	Ch	Yes	NO
4	No	Ch	Yes	NO
5	No	No	Yes	>30
6	No	Ch	Yes	NO
7	No	No	Yes	>30
8	No	Ch	Yes	NO
9	No	Ch	Yes	NO

[10 rows x 50 columns]

```
In [ ]: #Displaying variables data types  
print("Dataset Information:")  
print(df.info())
```

Dataset Information:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 101766 entries, 0 to 101765

Data columns (total 50 columns):

#	Column	Non-Null Count	Dtype
0	encounter_id	101766 non-null	int64
1	patient_nbr	101766 non-null	int64
2	race	101766 non-null	object
3	gender	101766 non-null	object
4	age	101766 non-null	object
5	weight	101766 non-null	object
6	admission_type_id	101766 non-null	int64
7	discharge_disposition_id	101766 non-null	int64
8	admission_source_id	101766 non-null	int64
9	time_in_hospital	101766 non-null	int64
10	payer_code	101766 non-null	object
11	medical_specialty	101766 non-null	object
12	num_lab_procedures	101766 non-null	int64
13	num_procedures	101766 non-null	int64
14	num_medications	101766 non-null	int64
15	number_outpatient	101766 non-null	int64
16	number_emergency	101766 non-null	int64
17	number_inpatient	101766 non-null	int64
18	diag_1	101766 non-null	object
19	diag_2	101766 non-null	object
20	diag_3	101766 non-null	object
21	number_diagnoses	101766 non-null	int64
22	max_glu_serum	5346 non-null	object
23	A1Cresult	17018 non-null	object
24	metformin	101766 non-null	object
25	repaglinide	101766 non-null	object
26	nateglinide	101766 non-null	object
27	chlorpropamide	101766 non-null	object
28	glimepiride	101766 non-null	object
29	acetohexamide	101766 non-null	object
30	glipizide	101766 non-null	object
31	glyburide	101766 non-null	object
32	tolbutamide	101766 non-null	object
33	pioglitazone	101766 non-null	object
34	rosiglitazone	101766 non-null	object
35	acarbose	101766 non-null	object
36	miglitol	101766 non-null	object
37	troglitazone	101766 non-null	object
38	tolazamide	101766 non-null	object
39	examide	101766 non-null	object
40	citoglipton	101766 non-null	object
41	insulin	101766 non-null	object
42	glyburide-metformin	101766 non-null	object
43	glipizide-metformin	101766 non-null	object
44	glimepiride-pioglitazone	101766 non-null	object
45	metformin-rosiglitazone	101766 non-null	object
46	metformin-pioglitazone	101766 non-null	object
47	change	101766 non-null	object
48	diabetesMed	101766 non-null	object
49	readmitted	101766 non-null	object

dtypes: int64(13), object(37)

memory usage: 38.8+ MB

None

```
In [ ]: #Defining lists for categorical variables and numeric variables
categorical_columns = [
    'race',
```

```

    'gender',
    'age',
    'weight',
    'payer_code',
    'medical_specialty',
    'diag_1',
    'diag_2',
    'diag_3',
    'max_glu_serum',
    'A1Cresult',
    'metformin',
    'repaglinide',
    'nateglinide',
    'chlorpropamide',
    'glimepiride',
    'acetoexamide',
    'glipizide',
    'glyburide',
    'tolbutamide',
    'pioglitazone',
    'rosiglitazone',
    'acarbose',
    'miglitol',
    'troglitazone',
    'tolazamide',
    'examide',
    'citoglipton',
    'insulin',
    'glyburide-metformin',
    'glipizide-metformin',
    'glimepiride-pioglitazone',
    'metformin-rosiglitazone',
    'metformin-pioglitazone',
    'change',
    'diabetesMed',
    'readmitted'
]

numeric_columns = [
    'encounter_id',
    'patient_nbr',
    'admission_type_id',
    'discharge_disposition_id',
    'admission_source_id',
    'time_in_hospital',
    'num_lab_procedures',
    'num_procedures',
    'num_medications',
    'number_outpatient',
    'number_emergency',
    'number_inpatient',
    'number_diagnoses'
]

```

```

In [ ]: #Dataset Description
print("Dataset Description:")
print(df.describe(include='all'))

```

## Dataset Description:

	encounter_id	patient_nbr	race	gender	age	weight	\
count	1.017660e+05	1.017660e+05	101766	101766	101766	101766	
unique	NaN	NaN	6	3	10	10	
top	NaN	NaN	Caucasian	Female	[70-80)	?	
freq	NaN	NaN	76099	54708	26068	98569	
mean	1.652016e+08	5.433040e+07	NaN	NaN	NaN	NaN	
std	1.026403e+08	3.869636e+07	NaN	NaN	NaN	NaN	
min	1.252200e+04	1.350000e+02	NaN	NaN	NaN	NaN	
25%	8.496119e+07	2.341322e+07	NaN	NaN	NaN	NaN	
50%	1.523890e+08	4.550514e+07	NaN	NaN	NaN	NaN	
75%	2.302709e+08	8.754595e+07	NaN	NaN	NaN	NaN	
max	4.438672e+08	1.895026e+08	NaN	NaN	NaN	NaN	

	admission_type_id	discharge_disposition_id	admission_source_id	\
count	101766.000000	101766.000000	101766.000000	
unique	NaN	NaN	NaN	
top	NaN	NaN	NaN	
freq	NaN	NaN	NaN	
mean	2.024006	3.715642	5.754437	
std	1.445403	5.280166	4.064081	
min	1.000000	1.000000	1.000000	
25%	1.000000	1.000000	1.000000	
50%	1.000000	1.000000	7.000000	
75%	3.000000	4.000000	7.000000	
max	8.000000	28.000000	25.000000	

	time_in_hospital	...	citoglipton	insulin	glyburide-metformin	\
count	101766.000000	...	101766	101766	101766	
unique	NaN	...	1	4	4	
top	NaN	...	No	No	No	
freq	NaN	...	101766	47383	101060	
mean	4.395987	...	NaN	NaN	NaN	
std	2.985108	...	NaN	NaN	NaN	
min	1.000000	...	NaN	NaN	NaN	
25%	2.000000	...	NaN	NaN	NaN	
50%	4.000000	...	NaN	NaN	NaN	
75%	6.000000	...	NaN	NaN	NaN	
max	14.000000	...	NaN	NaN	NaN	

	glipizide-metformin	glimepiride-pioglitazone	\
count	101766	101766	
unique	2	2	
top	No	No	
freq	101753	101765	
mean	NaN	NaN	
std	NaN	NaN	
min	NaN	NaN	
25%	NaN	NaN	
50%	NaN	NaN	
75%	NaN	NaN	
max	NaN	NaN	

	metformin-rosiglitazone	metformin-pioglitazone	change diabetesM	\
count	101766	101766	101766	101766
unique	2	2	2	2
top	No	No	No	Y
freq	101764	101765	54755	783

63				
mean	NaN	NaN	NaN	N
aN				
std	NaN	NaN	NaN	N
aN				
min	NaN	NaN	NaN	N
aN				
25%	NaN	NaN	NaN	N
aN				
50%	NaN	NaN	NaN	N
aN				
75%	NaN	NaN	NaN	N
aN				
max	NaN	NaN	NaN	N
aN				

	readmitted
count	101766
unique	3
top	NO
freq	54864
mean	NaN
std	NaN
min	NaN
25%	NaN
50%	NaN
75%	NaN
max	NaN

[11 rows x 50 columns]

```
In [ ]: #Converted blank and "?" to NaN to represent missing values
# Replace '?' and blank cells with NaN
df.replace({'?': pd.NA, '': pd.NA}, inplace=True)

# Replaced 'None' in 'max_glu_serum' and 'AlCresult' to be seen as 'No Test'
df['max_glu_serum'].replace({pd.NA: 'No Test', 'None': 'No Test'}, inplace=True)
df['AlCresult'].replace({pd.NA: 'No Test', 'None': 'No Test'}, inplace=True)
missing_values = df.isna().sum()
print("Number of missing values in each column:")
print(missing_values)
```



Number of missing values in each column:

encounter_id	0
patient_nbr	0
race	2273
gender	0
age	0
weight	98569
admission_type_id	0
discharge_disposition_id	0
admission_source_id	0
time_in_hospital	0
payer_code	40256
medical_specialty	49949
num_lab_procedures	0
num_procedures	0
num_medications	0
number_outpatient	0
number_emergency	0
number_inpatient	0
diag_1	21
diag_2	358
diag_3	1423
number_diagnoses	0
max_glu_serum	0
AlCresult	0
metformin	0
repaglinide	0
nateglinide	0
chlorpropamide	0
glimepiride	0
acetohexamide	0
glipizide	0
glyburide	0
tolbutamide	0
pioglitazone	0
rosiglitazone	0
acarbose	0
miglitol	0
troglitazone	0
tolazamide	0
examide	0
citoglipton	0
insulin	0
glyburide-metformin	0
glipizide-metformin	0
glimepiride-pioglitazone	0
metformin-rosiglitazone	0
metformin-pioglitazone	0
change	0
diabetesMed	0
readmitted	0

dtype: int64

```
In [ ]: # converting numeric columns to numeric data types and categorical columns
for col in numeric_columns:
    df[col] = pd.to_numeric(df[col], errors='coerce')

for col in categorical_columns:
    if col in df.columns:
        df[col] = df[col].astype('category')
print(df.dtypes)
print(df[numeric_columns].dtypes)
```

encounter_id	int64
patient_nbr	int64
race	category
gender	category
age	category
weight	category
admission_type_id	int64
discharge_disposition_id	int64
admission_source_id	int64
time_in_hospital	int64
payer_code	category
medical_specialty	category
num_lab_procedures	int64
num_procedures	int64
num_medications	int64
number_outpatient	int64
number_emergency	int64
number_inpatient	int64
diag_1	category
diag_2	category
diag_3	category
number_diagnoses	int64
max_glu_serum	category
AlCresult	category
metformin	category
repaglinide	category
nateglinide	category
chlorpropamide	category
glimepiride	category
acetohexamide	category
glipizide	category
glyburide	category
tolbutamide	category
pioglitazone	category
rosiglitazone	category
acarbose	category
miglitol	category
troglitazone	category
tolazamide	category
examide	category
citoglipton	category
insulin	category
glyburide-metformin	category
glipizide-metformin	category
glimepiride-pioglitazone	category
metformin-rosiglitazone	category
metformin-pioglitazone	category
change	category
diabetesMed	category
readmitted	category
dtype: object	
encounter_id	int64
patient_nbr	int64
admission_type_id	int64
discharge_disposition_id	int64
admission_source_id	int64
time_in_hospital	int64
num_lab_procedures	int64
num_procedures	int64
num_medications	int64
number_outpatient	int64
number_emergency	int64
number_inpatient	int64

```
number_diagnoses          int64
dtype: object
```

```
In [ ]: # Handle missing values by replacing with mode
        for column in ['race', 'diag_1', 'diag_2', 'diag_3']:
            df[column].fillna(df[column].mode()[0], inplace=True)

        # Verify changes
        print(df.isnull().sum())
```

```
encounter_id          0
patient_nbr           0
race                  0
gender                0
age                  0
weight              98569
admission_type_id     0
discharge_disposition_id 0
admission_source_id   0
time_in_hospital      0
payer_code            40256
medical_specialty     49949
num_lab_procedures    0
num_procedures         0
num_medications        0
number_outpatient      0
number_emergency       0
number_inpatient       0
diag_1                0
diag_2                0
diag_3                0
number_diagnoses      0
max_glu_serum         0
A1Cresult             0
metformin             0
repaglinide           0
nateglinide           0
chlorpropamide        0
glimepiride           0
acetohexamide         0
glipizide             0
glyburide             0
tolbutamide           0
pioglitazone          0
rosiglitazone         0
acarbose              0
miglitol              0
troglitazone          0
tolazamide            0
examide               0
citoglipton           0
insulin               0
glyburide-metformin   0
glipizide-metformin   0
glimepiride-pioglitazone 0
metformin-rosiglitazone 0
metformin-pioglitazone 0
change               0
diabetesMed           0
readmitted            0
dtype: int64
```

```
In [ ]: # Check for missing values
missing_values = df.isnull().sum()
print(missing_values)
```

```

encounter_id          0
patient_nbr           0
race                  0
gender                0
age                   0
weight                98569
admission_type_id     0
discharge_disposition_id 0
admission_source_id   0
time_in_hospital      0
payer_code            40256
medical_specialty     49949
num_lab_procedures    0
num_procedures        0
num_medications       0
number_outpatient      0
number_emergency       0
number_inpatient       0
diag_1                 0
diag_2                 0
diag_3                 0
number_diagnoses       0
max_glu_serum          0
AlCresult              0
metformin              0
repaglinide            0
nateglinide            0
chlorpropamide         0
glimepiride            0
acetohexamide          0
glipizide              0
glyburide              0
tolbutamide            0
pioglitazone           0
rosiglitazone          0
acarbose               0
miglitol               0
troglitazone           0
tolazamide             0
examide                0
citoglipton            0
insulin                0
glyburide-metformin    0
glipizide-metformin    0
glimepiride-pioglitazone 0
metformin-rosiglitazone 0
metformin-pioglitazone 0
change                 0
diabetesMed            0
readmitted             0
dtype: int64

```

```
In [ ]: #Checking for duplicate records
duplicate_records = df.duplicated().sum()
print(f"Number of duplicate records: {duplicate_records}")
```

```
Number of duplicate records: 0
```

```
In [ ]: # Drop unnecessary columns
df.drop(columns=['weight', 'payer_code', 'medical_specialty', 'encounter_

# Verify the remaining columns
print("Remaining columns after dropping unnecessary ones:")
print(df.columns)
```

```
Remaining columns after dropping unnecessary ones:
Index(['race', 'gender', 'age', 'admission_type_id',
       'discharge_disposition_id', 'admission_source_id', 'time_in_hospit
al',
       'num_lab_procedures', 'num_procedures', 'num_medications',
       'number_outpatient', 'number_emergency', 'number_inpatient', 'diag
_1',
       'diag_2', 'diag_3', 'number_diagnoses', 'max_glu_serum', 'A1Cresul
t',
       'metformin', 'repaglinide', 'nateglinide', 'chlorpropamide',
       'glimepiride', 'acetohehexamide', 'glipizide', 'glyburide', 'tolbuta
mide',
       'pioglitazone', 'rosiglitazone', 'acarbose', 'miglitol', 'troglita
zone',
       'tolazamide', 'examide', 'citoglipton', 'insulin',
       'glyburide-metformin', 'glipizide-metformin',
       'glimepiride-pioglitazone', 'metformin-rosiglitazone',
       'metformin-pioglitazone', 'change', 'diabetesMed', 'readmitted'],
      dtype='object')
```

```
In [ ]: # Simplify and categorize diag_1, diag_2, diag_3
def simplify_diag(x):
    if pd.isna(x):
        return -1
    if 'V' in x or 'E' in x:
        return 0
    try:
        return int(x[:3])
    except:
        return -1

df['diag_1'] = df['diag_1'].apply(simplify_diag)
df['diag_2'] = df['diag_2'].apply(simplify_diag)
df['diag_3'] = df['diag_3'].apply(simplify_diag)

# Verify changes
print(df[['diag_1', 'diag_2', 'diag_3']].head())
```

	diag_1	diag_2	diag_3
0	250	276	250
1	276	250	255
2	648	250	0
3	8	250	403
4	197	157	250

```
In [ ]: # Combine 'NO' and '>30' categories into a single '>30' category
df['readmitted'] = df['readmitted'].replace({'NO': '>30', '>30': '>30', '

# Verify changes
print(df['readmitted'].value_counts())
```

```
readmitted
>30    90409
<30    11357
Name: count, dtype: int64
```

```
In [ ]: # Here I dropped the List of columns dominated by "No" category since the
columns_dominated_by_no = [
    'metformin', 'repaglinide', 'nateglinide', 'chlorpropamide',
    'glimepiride', 'acetohexamide', 'glipizide', 'glyburide',
    'tolbutamide', 'pioglitazone', 'rosiglitazone', 'acarbose',
    'troglitazone', 'tolazamide', 'examide', 'citoglipton',
    'glyburide-metformin', 'glipizide-metformin', 'glimepiride-pioglitazo
    'metformin-rosiglitazone', 'metformin-pioglitazone'
]
df.drop(columns=columns_dominated_by_no, inplace=True)

# Check the remaining columns
print("Remaining columns after dropping unnecessary ones:")
print(df.columns)
print(df.info())
```

```
Remaining columns after dropping unnecessary ones:
Index(['race', 'gender', 'age', 'admission_type_id',
       'discharge_disposition_id', 'admission_source_id', 'time_in_hospit
al',
       'num_lab_procedures', 'num_procedures', 'num_medications',
       'number_outpatient', 'number_emergency', 'number_inpatient', 'diag
_1',
       'diag_2', 'diag_3', 'number_diagnoses', 'max_glu_serum', 'A1Cresul
t',
       'miglitol', 'insulin', 'change', 'diabetesMed', 'readmitted'],
      dtype='object')
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   race                                  101766 non-null  category
1   gender                                101766 non-null  category
2   age                                   101766 non-null  category
3   admission_type_id                     101766 non-null  int64
4   discharge_disposition_id              101766 non-null  int64
5   admission_source_id                   101766 non-null  int64
6   time_in_hospital                      101766 non-null  int64
7   num_lab_procedures                    101766 non-null  int64
8   num_procedures                        101766 non-null  int64
9   num_medications                       101766 non-null  int64
10  number_outpatient                     101766 non-null  int64
11  number_emergency                       101766 non-null  int64
12  number_inpatient                      101766 non-null  int64
13  diag_1                                101766 non-null  int64
14  diag_2                                101766 non-null  int64
15  diag_3                                101766 non-null  int64
16  number_diagnoses                      101766 non-null  int64
17  max_glu_serum                         101766 non-null  category
18  A1Cresult                             101766 non-null  category
19  miglitol                              101766 non-null  category
20  insulin                               101766 non-null  category
21  change                                101766 non-null  category
22  diabetesMed                           101766 non-null  category
23  readmitted                            101766 non-null  category
dtypes: category(10), int64(14)
memory usage: 11.8 MB
None
```

```
In [ ]: # Ensure that 'admission_type_id', 'discharge_disposition_id', 'admission
df[['admission_type_id', 'discharge_disposition_id', 'admission_source_id
```

```

# List of categorical features
categorical_features = ['race', 'gender', 'age', 'max_glu_serum', 'A1Cres',
                        'admission_type_id', 'discharge_disposition_id',

# Convert all categorical columns to strings
df[categorical_features] = df[categorical_features].astype(str)

# Verify conversion
print(df[categorical_features].dtypes)

```

```

race                object
gender              object
age                 object
max_glu_serum       object
A1Cresult            object
miglitol             object
insulin              object
change              object
diabetesMed          object
admission_type_id   object
discharge_disposition_id object
admission_source_id object
dtype: object

```

```

In [ ]: # One-hot encode categorical features
df = pd.get_dummies(df, columns=categorical_features, drop_first=True)

# Verify the changes
print(df.head())
print(df.info())

```

	time_in_hospital	num_lab_procedures	num_procedures	num_medications
0	1	41	0	1
1	3	59	0	18
2	2	11	5	13
3	2	44	1	16
4	1	51	0	8

	number_outpatient	number_emergency	number_inpatient	diag_1	diag_2
0	0	0	0	250	276
1	0	0	0	276	250
2	2	0	1	648	250
3	0	0	0	8	250
4	0	0	0	197	157

	diag_3	...	admission_source_id_20	admission_source_id_22	\
0	250	...	False	False	
1	255	...	False	False	
2	0	...	False	False	
3	403	...	False	False	
4	250	...	False	False	

	admission_source_id_25	admission_source_id_3	admission_source_id_4
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False

	admission_source_id_5	admission_source_id_6	admission_source_id_7	\
0	False	False	False	
1	False	False	True	
2	False	False	True	
3	False	False	True	
4	False	False	True	

	admission_source_id_8	admission_source_id_9
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False

[5 rows x 89 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 101766 entries, 0 to 101765

Data columns (total 89 columns):

#	Column	Non-Null Count	Dtype
0	time_in_hospital	101766 non-null	int64
1	num_lab_procedures	101766 non-null	int64
2	num_procedures	101766 non-null	int64
3	num_medications	101766 non-null	int64
4	number_outpatient	101766 non-null	int64
5	number_emergency	101766 non-null	int64
6	number_inpatient	101766 non-null	int64
7	diag_1	101766 non-null	int64
8	diag_2	101766 non-null	int64
9	diag_3	101766 non-null	int64
10	number_diagnoses	101766 non-null	int64
11	readmitted	101766 non-null	category



12	race_Asian	101766	non-null	bool
13	race_Caucasian	101766	non-null	bool
14	race_Hispanic	101766	non-null	bool
15	race_Other	101766	non-null	bool
16	gender_Male	101766	non-null	bool
17	gender_Unknown/Invalid	101766	non-null	bool
18	age_[10-20)	101766	non-null	bool
19	age_[20-30)	101766	non-null	bool
20	age_[30-40)	101766	non-null	bool
21	age_[40-50)	101766	non-null	bool
22	age_[50-60)	101766	non-null	bool
23	age_[60-70)	101766	non-null	bool
24	age_[70-80)	101766	non-null	bool
25	age_[80-90)	101766	non-null	bool
26	age_[90-100)	101766	non-null	bool
27	max_glu_serum_>300	101766	non-null	bool
28	max_glu_serum_No Test	101766	non-null	bool
29	max_glu_serum_Norm	101766	non-null	bool
30	AlCresult_>8	101766	non-null	bool
31	AlCresult_No Test	101766	non-null	bool
32	AlCresult_Norm	101766	non-null	bool
33	miglitol_No	101766	non-null	bool
34	miglitol_Steady	101766	non-null	bool
35	miglitol_Up	101766	non-null	bool
36	insulin_No	101766	non-null	bool
37	insulin_Steady	101766	non-null	bool
38	insulin_Up	101766	non-null	bool
39	change_No	101766	non-null	bool
40	diabetesMed_Yes	101766	non-null	bool
41	admission_type_id_2	101766	non-null	bool
42	admission_type_id_3	101766	non-null	bool
43	admission_type_id_4	101766	non-null	bool
44	admission_type_id_5	101766	non-null	bool
45	admission_type_id_6	101766	non-null	bool
46	admission_type_id_7	101766	non-null	bool
47	admission_type_id_8	101766	non-null	bool
48	discharge_disposition_id_10	101766	non-null	bool
49	discharge_disposition_id_11	101766	non-null	bool
50	discharge_disposition_id_12	101766	non-null	bool
51	discharge_disposition_id_13	101766	non-null	bool
52	discharge_disposition_id_14	101766	non-null	bool
53	discharge_disposition_id_15	101766	non-null	bool
54	discharge_disposition_id_16	101766	non-null	bool
55	discharge_disposition_id_17	101766	non-null	bool
56	discharge_disposition_id_18	101766	non-null	bool
57	discharge_disposition_id_19	101766	non-null	bool
58	discharge_disposition_id_2	101766	non-null	bool
59	discharge_disposition_id_20	101766	non-null	bool
60	discharge_disposition_id_22	101766	non-null	bool
61	discharge_disposition_id_23	101766	non-null	bool
62	discharge_disposition_id_24	101766	non-null	bool
63	discharge_disposition_id_25	101766	non-null	bool
64	discharge_disposition_id_27	101766	non-null	bool
65	discharge_disposition_id_28	101766	non-null	bool
66	discharge_disposition_id_3	101766	non-null	bool
67	discharge_disposition_id_4	101766	non-null	bool
68	discharge_disposition_id_5	101766	non-null	bool
69	discharge_disposition_id_6	101766	non-null	bool
70	discharge_disposition_id_7	101766	non-null	bool
71	discharge_disposition_id_8	101766	non-null	bool
72	discharge_disposition_id_9	101766	non-null	bool
73	admission_source_id_10	101766	non-null	bool
74	admission_source_id_11	101766	non-null	bool

```
75 admission_source_id_13      101766 non-null  bool
76 admission_source_id_14      101766 non-null  bool
77 admission_source_id_17      101766 non-null  bool
78 admission_source_id_2       101766 non-null  bool
79 admission_source_id_20      101766 non-null  bool
80 admission_source_id_22      101766 non-null  bool
81 admission_source_id_25      101766 non-null  bool
82 admission_source_id_3       101766 non-null  bool
83 admission_source_id_4       101766 non-null  bool
84 admission_source_id_5       101766 non-null  bool
85 admission_source_id_6       101766 non-null  bool
86 admission_source_id_7       101766 non-null  bool
87 admission_source_id_8       101766 non-null  bool
88 admission_source_id_9       101766 non-null  bool
dtypes: bool(77), category(1), int64(11)
memory usage: 16.1 MB
None
```

```
In [ ]: # Convert boolean columns to integer
bool_columns = df.select_dtypes(include='bool').columns
df[bool_columns] = df[bool_columns].astype(int)

# Verify the changes
print(df.head())
print(df.info())
```

	time_in_hospital	num_lab_procedures	num_procedures	num_medications
0	1	41	0	1
1	3	59	0	18
2	2	11	5	13
3	2	44	1	16
4	1	51	0	8

	number_outpatient	number_emergency	number_inpatient	diag_1	diag_2
0	0	0	0	250	276
1	0	0	0	276	250
2	2	0	1	648	250
3	0	0	0	8	250
4	0	0	0	197	157

	diag_3	...	admission_source_id_20	admission_source_id_22	\
0	250	...	0	0	
1	255	...	0	0	
2	0	...	0	0	
3	403	...	0	0	
4	250	...	0	0	

	admission_source_id_25	admission_source_id_3	admission_source_id_4
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

	admission_source_id_5	admission_source_id_6	admission_source_id_7	\
0	0	0	0	
1	0	0	1	
2	0	0	1	
3	0	0	1	
4	0	0	1	

	admission_source_id_8	admission_source_id_9
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

```
[5 rows x 89 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
Data columns (total 89 columns):
```

#	Column	Non-Null Count	Dtype
0	time_in_hospital	101766 non-null	int64
1	num_lab_procedures	101766 non-null	int64
2	num_procedures	101766 non-null	int64
3	num_medications	101766 non-null	int64
4	number_outpatient	101766 non-null	int64
5	number_emergency	101766 non-null	int64
6	number_inpatient	101766 non-null	int64
7	diag_1	101766 non-null	int64
8	diag_2	101766 non-null	int64
9	diag_3	101766 non-null	int64
10	number_diagnoses	101766 non-null	int64
11	readmitted	101766 non-null	category

12	race_Asian	101766	non-null	int64
13	race_Caucasian	101766	non-null	int64
14	race_Hispanic	101766	non-null	int64
15	race_Other	101766	non-null	int64
16	gender_Male	101766	non-null	int64
17	gender_Unknown/Invalid	101766	non-null	int64
18	age_[10-20)	101766	non-null	int64
19	age_[20-30)	101766	non-null	int64
20	age_[30-40)	101766	non-null	int64
21	age_[40-50)	101766	non-null	int64
22	age_[50-60)	101766	non-null	int64
23	age_[60-70)	101766	non-null	int64
24	age_[70-80)	101766	non-null	int64
25	age_[80-90)	101766	non-null	int64
26	age_[90-100)	101766	non-null	int64
27	max_glu_serum_>300	101766	non-null	int64
28	max_glu_serum_No Test	101766	non-null	int64
29	max_glu_serum_Norm	101766	non-null	int64
30	AlCresult_>8	101766	non-null	int64
31	AlCresult_No Test	101766	non-null	int64
32	AlCresult_Norm	101766	non-null	int64
33	miglitol_No	101766	non-null	int64
34	miglitol_Steady	101766	non-null	int64
35	miglitol_Up	101766	non-null	int64
36	insulin_No	101766	non-null	int64
37	insulin_Steady	101766	non-null	int64
38	insulin_Up	101766	non-null	int64
39	change_No	101766	non-null	int64
40	diabetesMed_Yes	101766	non-null	int64
41	admission_type_id_2	101766	non-null	int64
42	admission_type_id_3	101766	non-null	int64
43	admission_type_id_4	101766	non-null	int64
44	admission_type_id_5	101766	non-null	int64
45	admission_type_id_6	101766	non-null	int64
46	admission_type_id_7	101766	non-null	int64
47	admission_type_id_8	101766	non-null	int64
48	discharge_disposition_id_10	101766	non-null	int64
49	discharge_disposition_id_11	101766	non-null	int64
50	discharge_disposition_id_12	101766	non-null	int64
51	discharge_disposition_id_13	101766	non-null	int64
52	discharge_disposition_id_14	101766	non-null	int64
53	discharge_disposition_id_15	101766	non-null	int64
54	discharge_disposition_id_16	101766	non-null	int64
55	discharge_disposition_id_17	101766	non-null	int64
56	discharge_disposition_id_18	101766	non-null	int64
57	discharge_disposition_id_19	101766	non-null	int64
58	discharge_disposition_id_2	101766	non-null	int64
59	discharge_disposition_id_20	101766	non-null	int64
60	discharge_disposition_id_22	101766	non-null	int64
61	discharge_disposition_id_23	101766	non-null	int64
62	discharge_disposition_id_24	101766	non-null	int64
63	discharge_disposition_id_25	101766	non-null	int64
64	discharge_disposition_id_27	101766	non-null	int64
65	discharge_disposition_id_28	101766	non-null	int64
66	discharge_disposition_id_3	101766	non-null	int64
67	discharge_disposition_id_4	101766	non-null	int64
68	discharge_disposition_id_5	101766	non-null	int64
69	discharge_disposition_id_6	101766	non-null	int64
70	discharge_disposition_id_7	101766	non-null	int64
71	discharge_disposition_id_8	101766	non-null	int64
72	discharge_disposition_id_9	101766	non-null	int64
73	admission_source_id_10	101766	non-null	int64
74	admission_source_id_11	101766	non-null	int64

```

75 admission_source_id_13      101766 non-null   int64
76 admission_source_id_14      101766 non-null   int64
77 admission_source_id_17      101766 non-null   int64
78 admission_source_id_2       101766 non-null   int64
79 admission_source_id_20      101766 non-null   int64
80 admission_source_id_22      101766 non-null   int64
81 admission_source_id_25      101766 non-null   int64
82 admission_source_id_3       101766 non-null   int64
83 admission_source_id_4       101766 non-null   int64
84 admission_source_id_5       101766 non-null   int64
85 admission_source_id_6       101766 non-null   int64
86 admission_source_id_7       101766 non-null   int64
87 admission_source_id_8       101766 non-null   int64
88 admission_source_id_9       101766 non-null   int64
dtypes: category(1), int64(88)
memory usage: 68.4 MB
None

```

```

In [ ]: # Convert the target variable 'readmitted' into binary (1 for '<30', 0 for '>=30')
df['readmitted'] = df['readmitted'].apply(lambda x: 1 if x == '<30' else 0)

```

```

In [ ]: # To assess the normality of the numerical variables, I conducted the Shapiro-Wilk test.
# Despite the test results indicating that none of the variables followed a normal distribution,
# I chose to retain the outliers. The reason behind this decision was that the normality test
# across the numeric variables was not substantial, and removing them could have led to a loss of
# the variance and valuable information in the dataset.
from scipy.stats import shapiro

# Copy the DataFrame to avoid modifying the original
df_clustering = df.copy()

# List of numeric columns
numeric_columns = ['time_in_hospital', 'num_lab_procedures', 'num_procedures', 'number_outpatient',
                   'number_emergency', 'number_inpatient', 'diag_1', 'diag_2', 'diag_3']

# Dropping the class label 'readmitted' for Shapiro-Wilk test
data_for_shapiro = df_clustering[numeric_columns]

# Run Shapiro-Wilk test on numeric columns to check for normality
shapiro_results = {}
for column in numeric_columns:
    stat, p = shapiro(data_for_shapiro[column])
    shapiro_results[column] = {'statistic': stat, 'p_value': p}

# Display Shapiro-Wilk test results
print("Shapiro-Wilk Test Results:")
for column, result in shapiro_results.items():
    print(f"Column: {column}, Statistic: {result['statistic']}, p-value: {result['p_value']}")

```

```

Shapiro-Wilk Test Results:
Column: time_in_hospital, Statistic: 0.8869192600250244, p-value: 0.0
Column: num_lab_procedures, Statistic: 0.9848576188087463, p-value: 0.0
Column: num_procedures, Statistic: 0.7742846012115479, p-value: 0.0
Column: num_medications, Statistic: 0.9244745969772339, p-value: 0.0
Column: number_outpatient, Statistic: 0.3123285174369812, p-value: 0.0
Column: number_emergency, Statistic: 0.19770395755767822, p-value: 0.0
Column: number_inpatient, Statistic: 0.5557276010513306, p-value: 0.0
Column: diag_1, Statistic: 0.9683600068092346, p-value: 0.0
Column: diag_2, Statistic: 0.9584629535675049, p-value: 0.0
Column: diag_3, Statistic: 0.9462312459945679, p-value: 0.0

```

```
/usr/local/lib/python3.10/dist-packages/scipy/stats/_morestats.py:1882: UserWarning: p-value may not be accurate for N > 5000.
  warnings.warn("p-value may not be accurate for N > 5000.")
```

```
In [ ]: from sklearn.preprocessing import StandardScaler

# Normalize numeric columns
scaler = StandardScaler()
df_clustering[numeric_columns] = scaler.fit_transform(df_clustering[numeric_columns])

# Verify normalization
print("Normalized Data:")
print(pd.DataFrame(df_clustering[numeric_columns], columns=numeric_columns))
```

Normalized Data:

	time_in_hospital	num_lab_procedures	num_procedures	num_medications
count	1.017660e+05	1.017660e+05	1.017660e+05	1.017660e+05
mean	5.082986e-17	1.111554e-16	-1.717602e-17	-1.323811e-16
std	1.000005e+00	1.000005e+00	1.000005e+00	1.000005e+00
min	-1.137649e+00	-2.139630e+00	-7.853977e-01	-1.848268e+00
25%	-8.026506e-01	-6.147950e-01	-7.853977e-01	-7.409197e-01
50%	-1.326548e-01	4.596660e-02	-1.991621e-01	-1.257264e-01
75%	5.373411e-01	7.067282e-01	3.870736e-01	4.894670e-01
max	3.217324e+00	4.518815e+00	2.732016e+00	7.994826e+00

	number_outpatient	number_emergency	number_inpatient	diag_1
count	1.017660e+05	1.017660e+05	1.017660e+05	1.017660e+05
mean	2.010851e-17	3.044206e-17	-2.115583e-17	-1.273365e-16
std	1.000005e+00	1.000005e+00	1.000005e+00	1.000005e+00
min	-2.914615e-01	-2.126202e-01	-5.032762e-01	-2.265768e+00
25%	-2.914615e-01	-2.126202e-01	-5.032762e-01	-3.805559e-01
50%	-2.914615e-01	-2.126202e-01	-5.032762e-01	-2.358985e-01
75%	-2.914615e-01	-2.126202e-01	2.885790e-01	5.060539e-01
max	3.285094e+01	8.146673e+01	1.612568e+01	2.395933e+01


	diag_2	diag_3
count	1.017660e+05	1.017660e+05
mean	-3.239705e-17	-3.072134e-17
std	1.000005e+00	1.000005e+00
min	-2.255769e+00	-2.026482e+00
25%	-7.982336e-01	-7.439787e-01
50%	-1.665646e-02	3.065351e-02
75%	4.797506e-01	5.064624e-01
max	3.019876e+00	3.098402e+00

```
In [ ]: from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

# Prepare the data for clustering (excluding the 'readmitted' column)
clustering_data = df_clustering.drop(columns=['readmitted'])

# Calculate WCSS for different numbers of clusters
wcss = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=50)
    kmeans.fit(clustering_data)
    wcss.append(kmeans.inertia_)

# Visualize the Elbow Method
plt.figure(figsize=(10, 6))
plt.plot(range(1, 11), wcss, marker='o')
plt.title('Elbow Method for Optimal Number of Clusters')
plt.xlabel('Number of clusters (k)')
plt.ylabel('WCSS')
```



```
In [ ]: # Verify that clustering_data contains all the columns
print("Columns in clustering_data:")
print(clustering_data.columns)
# Get the count of columns
column_count = len(clustering_data.columns)
print(f"Number of columns in clustering_data: {column_count}")
```

```

Columns in clustering_data:
Index(['time_in_hospital', 'num_lab_procedures', 'num_procedures',
      'num_medications', 'number_outpatient', 'number_emergency',
      'number_inpatient', 'diag_1', 'diag_2', 'diag_3', 'number_diagnoses',
      'race_Asian', 'race_Caucasian', 'race_Hispanic', 'race_Other',
      'gender_Male', 'gender_Unknown/Invalid', 'age_[10-20)', 'age_[20-30)',
      'age_[30-40)', 'age_[40-50)', 'age_[50-60)', 'age_[60-70)',
      'age_[70-80)', 'age_[80-90)', 'age_[90-100)', 'max_glu_serum_>30',
      'max_glu_serum_No Test', 'max_glu_serum_Norm', 'A1Cresult_>8',
      'A1Cresult_No Test', 'A1Cresult_Norm', 'miglitol_No', 'miglitol_Steady',
      'miglitol_Up', 'insulin_No', 'insulin_Steady', 'insulin_Up',
      'change_No', 'diabetesMed_Yes', 'admission_type_id_2',
      'admission_type_id_3', 'admission_type_id_4', 'admission_type_id_5',
      'admission_type_id_6', 'admission_type_id_7', 'admission_type_id_8',
      'discharge_disposition_id_10', 'discharge_disposition_id_11',
      'discharge_disposition_id_12', 'discharge_disposition_id_13',
      'discharge_disposition_id_14', 'discharge_disposition_id_15',
      'discharge_disposition_id_16', 'discharge_disposition_id_17',
      'discharge_disposition_id_18', 'discharge_disposition_id_19',
      'discharge_disposition_id_2', 'discharge_disposition_id_20',
      'discharge_disposition_id_22', 'discharge_disposition_id_23',
      'discharge_disposition_id_24', 'discharge_disposition_id_25',
      'discharge_disposition_id_27', 'discharge_disposition_id_28',
      'discharge_disposition_id_3', 'discharge_disposition_id_4',
      'discharge_disposition_id_5', 'discharge_disposition_id_6',
      'discharge_disposition_id_7', 'discharge_disposition_id_8',
      'discharge_disposition_id_9', 'admission_source_id_10',
      'admission_source_id_11', 'admission_source_id_13',
      'admission_source_id_14', 'admission_source_id_17',
      'admission_source_id_2', 'admission_source_id_20',
      'admission_source_id_22', 'admission_source_id_25',
      'admission_source_id_3', 'admission_source_id_4',
      'admission_source_id_5', 'admission_source_id_6',
      'admission_source_id_7', 'admission_source_id_8',
      'admission_source_id_9'],
      dtype='object')
Number of columns in clustering_data: 88

```

```

In [ ]: from sklearn.cluster import KMeans
        from sklearn.metrics import silhouette_score
        import pandas as pd

        # Perform K-Means clustering with the optimal number of clusters (K=4)
        optimal_clusters = 4 # Based on the Elbow Method
        kmeans = KMeans(n_clusters=optimal_clusters, n_init=50, random_state=42)
        kmeans.fit(clustering_data)

        # Add the cluster labels to the DataFrame
        df_clustering['Cluster'] = kmeans.labels_

        # Print the DataFrame after adding the cluster labels
        print(df_clustering)

        # Calculate the silhouette score to evaluate the clustering performance
        silhouette_avg = silhouette_score(clustering_data, kmeans.labels_)
        print(f'Silhouette Score: {silhouette_avg}')

```



	time_in_hospital	num_lab_procedures	num_procedures	num_medication
0	-1.137649	-0.106517	-0.785398	-1.84
8268				
1	-0.467653	0.808384	-0.785398	0.24
3390				
2	-0.802651	-1.631351	2.145781	-0.37
1804				
3	-0.802651	0.045967	-0.199162	-0.00
2688				
4	-1.137649	0.401761	-0.785398	-0.98
6997				
...	...	...	...	
...				
101761	-0.467653	0.401761	-0.785398	-0.00
2688				
101762	0.202343	-0.513139	0.973309	0.24
3390				
101763	-1.137649	0.503417	-0.785398	-0.86
3958				
101764	1.877333	0.096794	0.387074	0.61
2506				
101765	0.537341	-1.529696	0.973309	-1.60
2190				

	number_outpatient	number_emergency	number_inpatient	diag_1
0	-0.291461	-0.21262	-0.503276	-1.099176
1	-0.291461	-0.21262	-0.503276	-0.977851
2	1.286748	-0.21262	0.288579	0.758038
3	-0.291461	-0.21262	-0.503276	-2.228437
4	-0.291461	-0.21262	-0.503276	-1.346494
...	...	...	...	...
...				
101761	-0.291461	-0.21262	-0.503276	-1.099176
101762	-0.291461	-0.21262	0.288579	0.347397
101763	0.497643	-0.21262	-0.503276	-2.088446
101764	-0.291461	-0.21262	0.288579	2.381934
101765	-0.291461	-0.21262	-0.503276	0.207406

	diag_2	diag_3	...	admission_source_id_22	\
0	-0.798234	-0.743979	...	0	
1	-0.935538	-0.718329	...	0	
2	-0.935538	-2.026482	...	0	
3	-0.935538	0.040914	...	0	
4	-1.426664	-0.743979	...	0	
...	...	...	...	...	
...					
101761	-0.719020	0.323064	...	0	
101762	-0.798234	2.010839	...	0	
101763	0.859977	-0.507998	...	0	
101764	-0.750705	3.093272	...	0	
101765	0.543122	2.010839	...	0	

	admission_source_id_25	admission_source_id_3	admission_source_id
4 \			
0	0	0	
0			
1	0	0	
0			
2	0	0	
0			
3	0	0	
0			
4	0	0	

```

0
...
...
101761      0      0
0
101762      0      0
0
101763      0      0
0
101764      0      0
0
101765      0      0
0

      admission_source_id_5  admission_source_id_6  admission_source_id
_7 \
0      0      0
0
1      0      0
1
2      0      0
1
3      0      0
1
4      0      0
1
...      ...      ...
...
101761      0      0
1
101762      1      0
0
101763      0      0
1
101764      0      0
1
101765      0      0
1

      admission_source_id_8  admission_source_id_9  Cluster
0      0      0      2
1      0      0      3
2      0      0      2
3      0      0      3
4      0      0      2
...      ...      ...
101761      0      0      3
101762      0      0      3
101763      0      0      3
101764      0      0      1
101765      0      0      3

```

```

[101766 rows x 90 columns]
Silhouette Score: 0.12615024043888765

```

```

In [ ]: from sklearn.cluster import KMeans
        from sklearn.metrics import silhouette_score
        import matplotlib.pyplot as plt

        # Function to evaluate and plot silhouette scores for different k values
        def evaluate_k_means(data, k_values, n_init=50, random_state=42):
            silhouette_scores = []
            for k in k_values:

```

```

kmeans = KMeans(n_clusters=k, random_state=random_state, n_init=n
kmeans.fit(data)
score = silhouette_score(data, kmeans.labels_)
silhouette_scores.append(score)
print(f'For n_clusters = {k}, Silhouette Score = {score}')

plt.figure(figsize=(10, 6))
plt.plot(k_values, silhouette_scores, marker='o')
plt.title('Silhouette Scores for Different k Values')
plt.xlabel('Number of clusters (k)')
plt.ylabel('Silhouette Score')
plt.show()

# Define range of k values to evaluate
k_values = range(2, 6)


# Evaluate k-means clustering
evaluate_k_means(clustering_data, k_values)

```

```

For n_clusters = 2, Silhouette Score = 0.18104549510866558
For n_clusters = 3, Silhouette Score = 0.12956704345897635
For n_clusters = 4, Silhouette Score = 0.12615024043888765
For n_clusters = 5, Silhouette Score = 0.13011944653549468

```



```

In [ ]: # Get the cluster centers
centers = kmeans.cluster_centers_
centers_df = pd.DataFrame(centers, columns=clustering_data.columns)

print("\nCluster Centers:")
print(centers_df)

```

## Cluster Centers:

	time_in_hospital	num_lab_procedures	num_procedures	num_medications
\				
0	-0.124821	0.074534	-0.283150	-0.283080
1	-0.493989	-1.084364	-0.405609	-0.419612
2	0.101573	0.653746	-0.439452	0.037352
3	1.747294	0.868241	1.087644	1.752257
4	-0.307474	-0.224819	1.877981	0.189714
5	-0.565761	-0.485532	-0.080871	-0.558479
6	-0.020726	0.059744	-0.171299	0.187226
7	0.189131	0.184523	-0.246129	0.199210
8	-0.002749	-0.014285	-0.224957	-0.113821
9	-0.148754	-0.061326	-0.288018	-0.023839

	number_outpatient	number_emergency	number_inpatient	diag_1	dia
g_2 \					
0	-0.161802	-0.118401	-0.185763	-0.000730	-0.024
152					
1	-0.057230	-0.043332	-0.176297	0.312814	-0.051
197					
2	-0.105205	-0.047106	-0.145427	-0.108873	-0.031
963					
3	-0.114240	-0.081855	-0.096228	-0.190869	0.263
911					
4	-0.134249	-0.113021	-0.260529	-0.031398	0.133
717					
5	-0.152540	-0.107597	-0.258944	0.037006	-0.481
932					
6	4.866870	0.222834	0.260631	-0.048029	0.135
032					
7	0.125006	0.470208	3.055249	-0.118085	0.032
696					
8	-0.085547	-0.050226	-0.137136	0.007271	0.449
355					
9	0.440758	8.883833	2.744648	-0.126580	-0.071
193					

	diag_3	...	admission_source_id_20	admission_source_id_22	\
0	-0.154916	...	9.067481e-04	4.772358e-05	
1	-0.325354	...	1.719937e-03	1.433281e-04	
2	-0.256922	...	1.813659e-03	2.072754e-04	
3	0.261574	...	3.039195e-03	2.095997e-04	
4	-0.024582	...	4.113111e-03	1.285347e-04	
5	-0.436995	...	3.773015e-04	3.876023e-18	
6	0.163059	...	-1.691355e-17	1.138412e-18	
7	0.033645	...	8.259343e-04	2.064836e-04	
8	1.661158	...	1.371742e-03	1.055186e-04	
9	-0.040067	...	1.734723e-18	-5.827587e-19	

	admission_source_id_25	admission_source_id_3	admission_source_id_4
\			
0	8.131516e-20	2.147561e-03	0.028682
1	2.093865e-18	9.316325e-04	0.021929
2	-5.454892e-19	8.809203e-04	0.016271
3	1.047998e-04	2.515196e-03	0.060050
4	8.402567e-19	2.442159e-03	0.077378
5	3.184844e-19	3.697555e-03	0.034033
6	-2.473336e-19	1.149254e-17	0.007529
7	-3.828589e-19	1.032418e-03	0.020442
8	1.055186e-04	1.582779e-03	0.023531
9	1.084202e-19	9.107298e-18	0.003697

admission_source_id_5	admission_source_id_6	admission_source_id_7	\
-----------------------	-----------------------	-----------------------	---

0	0.006538	6.762432e-02	0.540326
1	0.004443	3.869858e-03	0.505661
2	0.013473	7.410094e-03	0.755726
3	0.013624	7.440788e-03	0.521065
4	0.005784	3.084833e-03	0.389846
5	0.002188	3.471174e-02	0.431407
6	0.004235	6.117647e-03	0.573647
7	0.012183	7.020442e-03	0.697089
8	0.012873	5.064894e-03	0.608737
9	0.003697	1.387779e-17	0.824399

	admission_source_id_8	admission_source_id_9
0	2.386179e-04	3.531545e-03
1	7.166404e-05	2.866562e-04
2	5.181884e-05	-2.474149e-16
3	2.095997e-04	1.047998e-04
4	6.722053e-18	-2.905662e-17
5	3.773015e-04	3.395714e-03
6	-1.978669e-18	2.797242e-17
7	-3.062871e-18	3.989864e-17
8	1.055186e-04	1.055186e-04
9	1.848429e-03	-4.987330e-18

[10 rows x 88 columns]

```
In [ ]: from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import pandas as pd
# Perform K-Means clustering with the optimal number of clusters (K=2) si
optimal_clusters = 2
kmeans = KMeans(n_clusters=optimal_clusters, n_init=50, random_state=42)
kmeans.fit(clustering_data)

# Add the cluster labels to the clustering_data DataFrame
clustering_data['Cluster'] = kmeans.labels_

# Calculate the mean values for each feature (including one-hot encoded c
cluster_means = clustering_data.groupby('Cluster').mean()

print("\nMean values for each feature within each cluster:")
print(cluster_means)

# Calculate the silhouette score
silhouette_avg = silhouette_score(clustering_data.drop(columns=['Cluster'])
print(f"\nSilhouette Score for K={optimal_clusters}: {silhouette_avg}")

# Optional: Save the clustering results
clustering_data.to_csv('clustering_results.csv', index=False)
```

Mean values for each feature within each cluster:

	time_in_hospital	num_lab_procedures	num_procedures	\
Cluster				
0	-0.332139	-0.236397	-0.115492	
1	0.178380	0.126960	0.062027	

	num_medications	number_outpatient	number_emergency	\
Cluster				
0	-0.386925	-0.126819	-0.079964	
1	0.207804	0.068110	0.042946	

	number_inpatient	diag_1	diag_2	diag_3	...	\
Cluster					...	
0	-0.150618	0.020951	-0.215357	-0.264754	...	
1	0.080892	-0.011252	0.115661	0.142190	...	

	admission_source_id_20	admission_source_id_22	\
Cluster			
0	0.000562	0.000028	
1	0.002130	0.000166	

	admission_source_id_25	admission_source_id_3	admission_source_id_4	\
Cluster				
0	0.000000	0.002840	0.03	
3185				
1	0.000003	0.001299	0.03	
0314				

	admission_source_id_5	admission_source_id_6	admission_source_id_7	\
Cluster				
0	0.003712	0.053490	0.478	
570				
1	0.010920	0.005468	0.611	
361				

	admission_source_id_8	admission_source_id_9
Cluster		
0	0.000281	0.003347
1	0.000091	0.000091

[2 rows x 88 columns]

Silhouette Score for K=2: 0.18104549510866558

```
In [ ]: from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import pandas as pd

# Perform K-Means clustering with the optimal number of clusters (K=4)
optimal_clusters = 4 # Based on the Elbow Method
kmeans = KMeans(n_clusters=optimal_clusters, n_init=50, random_state=42)
kmeans.fit(clustering_data)

# Add the cluster labels to the clustering_data DataFrame
clustering_data['Cluster'] = kmeans.labels_

# Calculate the mean values for each feature (including one-hot encoded c
cluster_means = clustering_data.groupby('Cluster').mean()

print("\nMean values for each feature within each cluster:")
print(cluster_means)
```

```
# Calculate the silhouette score
silhouette_avg = silhouette_score(clustering_data.drop(columns=['Cluster'])
print(f"\nSilhouette Score for K={optimal_clusters}: {silhouette_avg}")

# Optional: Save the clustering results
clustering_data.to_csv('clustering_results.csv', index=False)
```

Mean values for each feature within each cluster:

	time_in_hospital	num_lab_procedures	num_procedures	\
Cluster				
0	1.188584	0.702316	0.788902	
1	-0.339837	-0.196827	-0.262453	
2	-0.319039	-0.208386	-0.108835	
3	0.055019	0.116061	-0.248414	

	num_medications	number_outpatient	number_emergency	\
Cluster				
0	1.106119	-0.097712	-0.085151	
1	-0.274204	-0.071718	-0.079213	
2	-0.387190	-0.144489	-0.115751	
3	0.213187	1.927709	1.770374	

	number_inpatient	diag_1	diag_2	diag_3	...	\
Cluster					...	
0	-0.052875	-0.184521	0.245164	0.308356	...	
1	-0.123027	0.089893	0.020994	0.017894	...	
2	-0.189744	0.009810	-0.211165	-0.255412	...	
3	2.479069	-0.106375	0.059963	0.091354	...	

	admission_source_id_20	admission_source_id_22	\
Cluster			
0	0.002773	0.000382	
1	0.001865	0.000044	
2	0.000562	0.000033	
3	0.000199	0.000199	

	admission_source_id_25	admission_source_id_3	admission_source_id_4	\
Cluster				
0	0.000096	0.002056	0.05	
1630				
1	0.000000	0.001163	0.02	
2838				
2	0.000000	0.002910	0.03	
3395				
3	0.000000	0.000597	0.01	
1149				

	admission_source_id_5	admission_source_id_6	admission_source_id_7	\
Cluster				
0	0.013386	0.006693	0.562	
673				
1	0.009434	0.005046	0.617	
670				
2	0.003505	0.061830	0.466	
373				
3	0.007764	0.004778	0.689	
827				

	admission_source_id_8	admission_source_id_9
Cluster		
0	0.000143	0.000143
1	0.000110	0.000066
2	0.000231	0.003935
3	0.000199	0.000000

[4 rows x 88 columns]

Silhouette Score for K=4: 0.1254200111853815



```
In [ ]: # Calculate the mean values for each feature (including one-hot encoded c
cluster_means = clustering_data.groupby('Cluster').mean()

print("\nMean values for each feature within each cluster:")
display(cluster_means)

clustering_data.to_csv('clustering_results.csv', index=False)
```

Mean values for each feature within each cluster:

	time_in_hospital	num_lab_procedures	num_procedures	num_medications	number_c
Cluster					
0	1.188584	0.702316	0.788902	1.106119	
1	-0.339837	-0.196827	-0.262453	-0.274204	
2	-0.319039	-0.208386	-0.108835	-0.387190	
3	0.055019	0.116061	-0.248414	0.213187	

4 rows × 88 columns

Warning: Total number of columns (88) exceeds max\_columns (20) limiting to first (20) columns.

```
In [ ]: import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

# Ensure the Cluster column is present in clustering_data
clustering_data['Cluster'] = kmeans.labels_

# Visualize the distribution of 'time_in_hospital' by cluster
plt.figure(figsize=(10, 6))
sns.histplot(data=clustering_data, x='time_in_hospital', hue='Cluster', e
plt.title('Distribution of Time in Hospital by Cluster')
plt.show()

# Visualize the distribution of 'num_lab_procedures' by cluster
plt.figure(figsize=(10, 6))
sns.histplot(data=clustering_data, x='num_lab_procedures', hue='Cluster',
plt.title('Distribution of Number of Lab Procedures by Cluster')
plt.show()

# Visualize the distribution of 'num_medications' by cluster
plt.figure(figsize=(10, 6))
sns.histplot(data=clustering_data, x='num_medications', hue='Cluster', el
plt.title('Distribution of Number of Medications by Cluster')
plt.show()

# Sum the one-hot encoded 'race' columns by cluster
race_columns = [col for col in clustering_data.columns if col.startswith(
race_distribution = clustering_data.groupby('Cluster')[race_columns].sum(

# Convert to long format for easier plotting
race_distribution = race_distribution.reset_index().melt(id_vars='Cluster


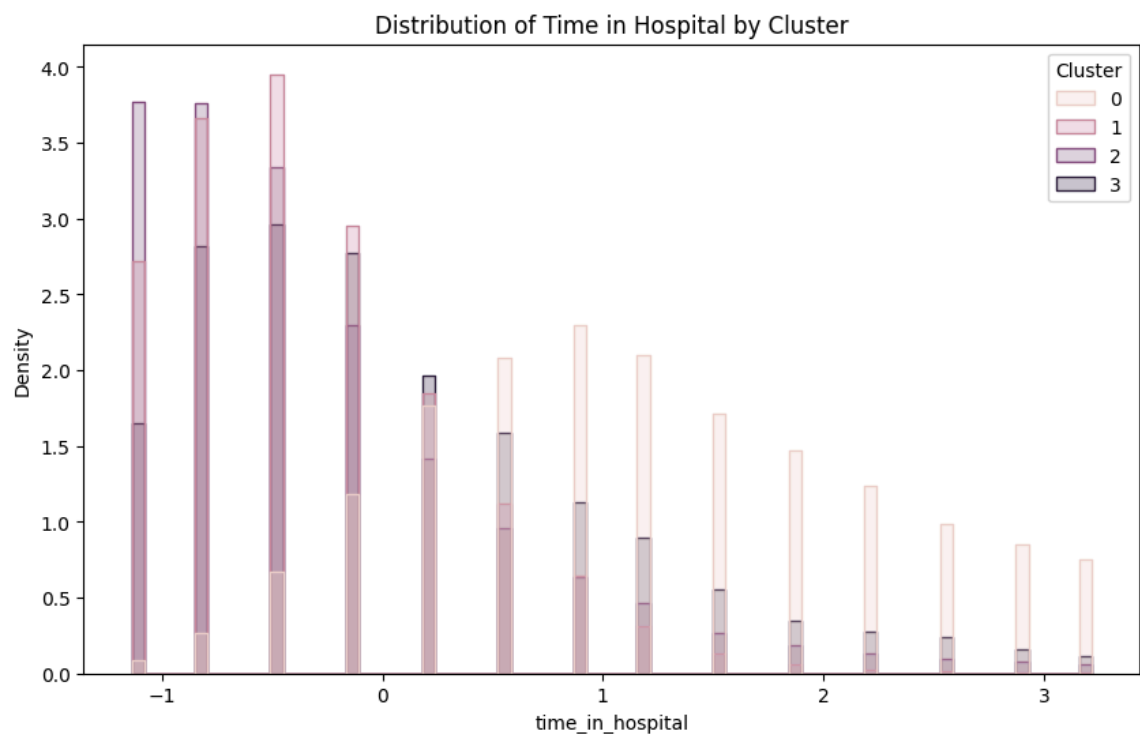
# Visualize the distribution of 'race' by cluster
plt.figure(figsize=(10, 6))
sns.barplot(data=race_distribution, x='Race', y='Count', hue='Cluster')
plt.title('Distribution of Race by Cluster')
plt.xticks(rotation=90)
```

```
plt.show()


# Sum the one-hot encoded 'age' columns by cluster
age_columns = [col for col in clustering_data.columns if col.startswith('age_')]
age_distribution = clustering_data.groupby('Cluster')[age_columns].sum()

# Convert to long format for easier plotting
age_distribution = age_distribution.reset_index().melt(id_vars='Cluster',


# Visualize the distribution of 'age' by cluster
plt.figure(figsize=(10, 6))
sns.barplot(data=age_distribution, x='Age', y='Count', hue='Cluster')
plt.title('Distribution of Age by Cluster')
plt.xticks(rotation=90)
plt.show()
```




Distribution of Number of Medications by Cluster



Distribution of Race by Cluster





```
In [ ]: from sklearn.cluster import KMeans
import seaborn as sns
import matplotlib.pyplot as plt

# Perform K-Means clustering with K=2
kmeans = KMeans(n_clusters=2, n_init=50, random_state=42)
kmeans.fit(clustering_data)

# Add the cluster labels to the DataFrame
clustering_data['Cluster'] = kmeans.labels_

# Visualize the distribution of 'time_in_hospital' by cluster
plt.figure(figsize=(10, 6))
sns.histplot(data=clustering_data, x='time_in_hospital', hue='Cluster', e
plt.title('Distribution of Time in Hospital by Cluster')
plt.show()

# Visualize the distribution of 'num_lab_procedures' by cluster
plt.figure(figsize=(10, 6))
sns.histplot(data=clustering_data, x='num_lab_procedures', hue='Cluster',
plt.title('Distribution of Number of Lab Procedures by Cluster')
plt.show()

# Visualize the distribution of 'num_medications' by cluster
plt.figure(figsize=(10, 6))
sns.histplot(data=clustering_data, x='num_medications', hue='Cluster', el
plt.title('Distribution of Number of Medications by Cluster')
plt.show()

# Visualize the distribution of 'race' by cluster
plt.figure(figsize=(10, 6))
race_columns = ['race_Asian', 'race_Caucasian', 'race_Hispanic', 'race_Ot
race_data = clustering_data.melt(id_vars='Cluster', value_vars=race_colum
race_data = race_data[race_data['Count'] == 1]
sns.countplot(data=race_data, x='Race', hue='Cluster')
plt.title('Distribution of Race by Cluster')
```

```
plt.show()

# Visualize the distribution of 'age' by cluster
plt.figure(figsize=(10, 6))
age_columns = ['age_[10-20)', 'age_[20-30)', 'age_[30-40)', 'age_[40-50)']
age_data = clustering_data.melt(id_vars='Cluster', value_vars=age_columns)
age_data = age_data[age_data['Count'] == 1]
sns.countplot(data=age_data, x='Age', hue='Cluster')
plt.title('Distribution of Age by Cluster')
plt.show()
```

